

Problem Solving-Blind Search

Instructor: Dr. Karen Hand

AI a Modern Approach. Chapter 3.



Russell and Norvig:
Chapter 3, Sections 3.1 – 3.2

Solving Problems by Searching

- AI main purpose is to solve problems.
- The first problems that AI had to solve involved searching.
- The problem-solving approach has been applied to a vast array of task environments.
- We list some of the best known here, distinguishing between toy and real-world problems.

Toy problem


- Is intended to illustrate or exercise various problem-solving methods. It can be given a concise, exact description and hence is usable by different researchers to compare the performance of algorithms.



Real-world problem

- Is one whose solutions people actually care about. Such problems tend not to have a single agreed-upon description, but we can give the general flavor of their formulations.

Problem-Solving Agent



Problem Solving as Search


For a goal-based agent, what should it do when none of the actions it can currently perform results in a goal state?

One way to address the issues is to view goal-attainment as **problem solving**, and viewing that as a **search** through the space of possible solutions.

Search Problem Formulation


- Real-world environment → Abstraction
 - States within the solver are abstractions of states of the world the agent is actually in
 - Actions in the search space are abstractions of the agent's actions
 - Solutions map to sequences of real actions
 - The path cost function that assigns a numeric cost to action sequence.
- Without abstraction an agent would be swamped by the real world

Problem-Solving Agent




- Actions
- Initial state
- Goal test
- Path Cost Function

State Space and Successor Function



- Actions
- Initial state
- Goal test
- Path Cost Function

Initial State




- Actions
- Initial state
- Goal test
- Path Cost Function

Goal Test

state space

successor function

- Actions
- Initial state
- Goal test
- Path Cost Function



Example: 8- puzzle

8	2	
3	4	7
5	1	6

Initial state

1	2	3
4	5	6
7	8	

Goal state


Example: 8-
puzzle

8	2	
3	4	7
5	1	6

8	2	7
3	4	
5	1	6

8		2
3	4	7
5	1	6

	8	2
3	4	7
5	1	6



Example: 8- puzzle

Size of the state space = $9!/2 = 181,440$

15-puzzle $\rightarrow 6.5 \times 10^{11}$

24-puzzle $\rightarrow 5 \times 10^{24}$

0.018 sec

18 hours

16 billion years

10 millions states/sec

Search Problem

- State space
- Initial state
- Successor function
- Goal test
- Path cost


What is a Solution?

- A sequence of actions that when performed will transform the initial state into a goal state
- Or sometimes just the goal state itself, when getting there is trivial


Assumptions in Basic Search

- The environment is **static**
- The environment is **discretizable**
- The environment is **observable**
- The actions are **deterministic**


Search of State Space




Search of State Space




Search of State Space




Search of State Space



Search of State Space



Search of State Space



→ search tree

Implementation : states vs. Nodes

- A state is (representation of) a physical configuration.
- A node is a data structure constituting part of a search tree includes **parent, children, depth, path cost**.
- States do not have parents, children, depth, or path cost.
- The expanding process creates new nodes, filling in the various fields and using the successor function of the problem to create the corresponding states.


Simple Agent Algorithm

Problem-Solving-Agent


1. $\text{initial-state} \leftarrow \text{sense/read state}$
2. $\text{goal} \leftarrow \text{select/read goal}$
3. $\text{successor} \leftarrow \text{select/read action models}$
4. $\text{problem} \leftarrow (\text{initial-state}, \text{goal}, \text{successor})$
5. $\text{solution} \leftarrow \text{search}(\text{problem})$
6. $\text{perform}(\text{solution})$

Example: 8-queens

Place 8 queens in a chessboard so that no two queens are in the same row, column, or diagonal.




A solution



Not a solution

Example: 8-queens




Formulation #1:

- States: any arrangement of 0 to 8 queens on the board
- Initial state: 0 queens on the board
- Successor function: add a queen in any square
- Goal test: 8 queens on the board, none attacked

→ 64^8 states with 8 queens

Example: 8-queens



→ 2,067 states

Formulation #2:

- States: any arrangement of $k = 0$ to 8 queens in the k leftmost columns with none attacked
- Initial state: 0 queens on the board
- Successor function: add a queen to any square in the leftmost empty column such that it is not attacked by any other queen
- Goal test: 8 queens on the board

Search Problem Variants

- One or several initial states
- One or several goal states
- The solution is the path or a goal node
 - In the 8-puzzle problem, it is the path to a goal node
 - In the 8-queen problem, it is a goal node
- Any, or the best, or all solutions

Important Parameters

8-puzzle \rightarrow 181,440

15-puzzle $\rightarrow .65 \times 10^{12}$

24-puzzle $\rightarrow .5 \times 10^{25}$

- Number of states in state space

8-queens \rightarrow 2,057

100-queens $\rightarrow 10^{52}$

There exist techniques to solve
N-queens problems efficiently!

**Stating a problem as a search problem
is not always a good idea!**

Important Parameters

- Number of states in state space
- Size of memory needed to store a state
- Running time of the successor function

Applications

- Route finding: airline travel, telephone/computer networks
- Touring problems
- VLSI layout
- Automatic assembly sequencing
- Internet searching

Summary

- Problem-solving agent
- State space, successor function, search
- Examples: 8-puzzle, 8-queens, route finding, robot navigation, assembly planning
- Assumptions of basic search
- Important parameters


Acknowledgement

- Krishnaram Kenthapadi , CS121: Introduction to Artificial Intelligence, 2003.
- Jim Martin, CSCI 5582: Introduction to Artificial Intelligence, 2003.
- **AI Instructor's Resource Page,**
<http://aima.cs.berkeley.edu/instructors.html>

Blind Search

Russell and Norvig:
Chapter 3, Sections 3.3 – 3.4

Robot Navigation




Simple Agent Algorithm

Problem-Solving-Agent

1. $\text{initial-state} \leftarrow \text{sense/read state}$
2. $\text{goal} \leftarrow \text{select/read goal}$
3. $\text{successor} \leftarrow \text{select/read action models}$
4. $\text{problem} \leftarrow (\text{initial-state}, \text{goal}, \text{successor})$
5. $\text{solution} \leftarrow \text{search}(\text{problem})$
6. $\text{perform}(\text{solution})$

Search of State Space



→ search tree

Basic Search Concepts

- Search tree
- Search node
- Node expansion
- **Search strategy:** At each stage it determines which node to expand

Search Nodes \neq
States

8	2	
3	4	7
5	1	6

8	2	7
3	4	
5	1	6

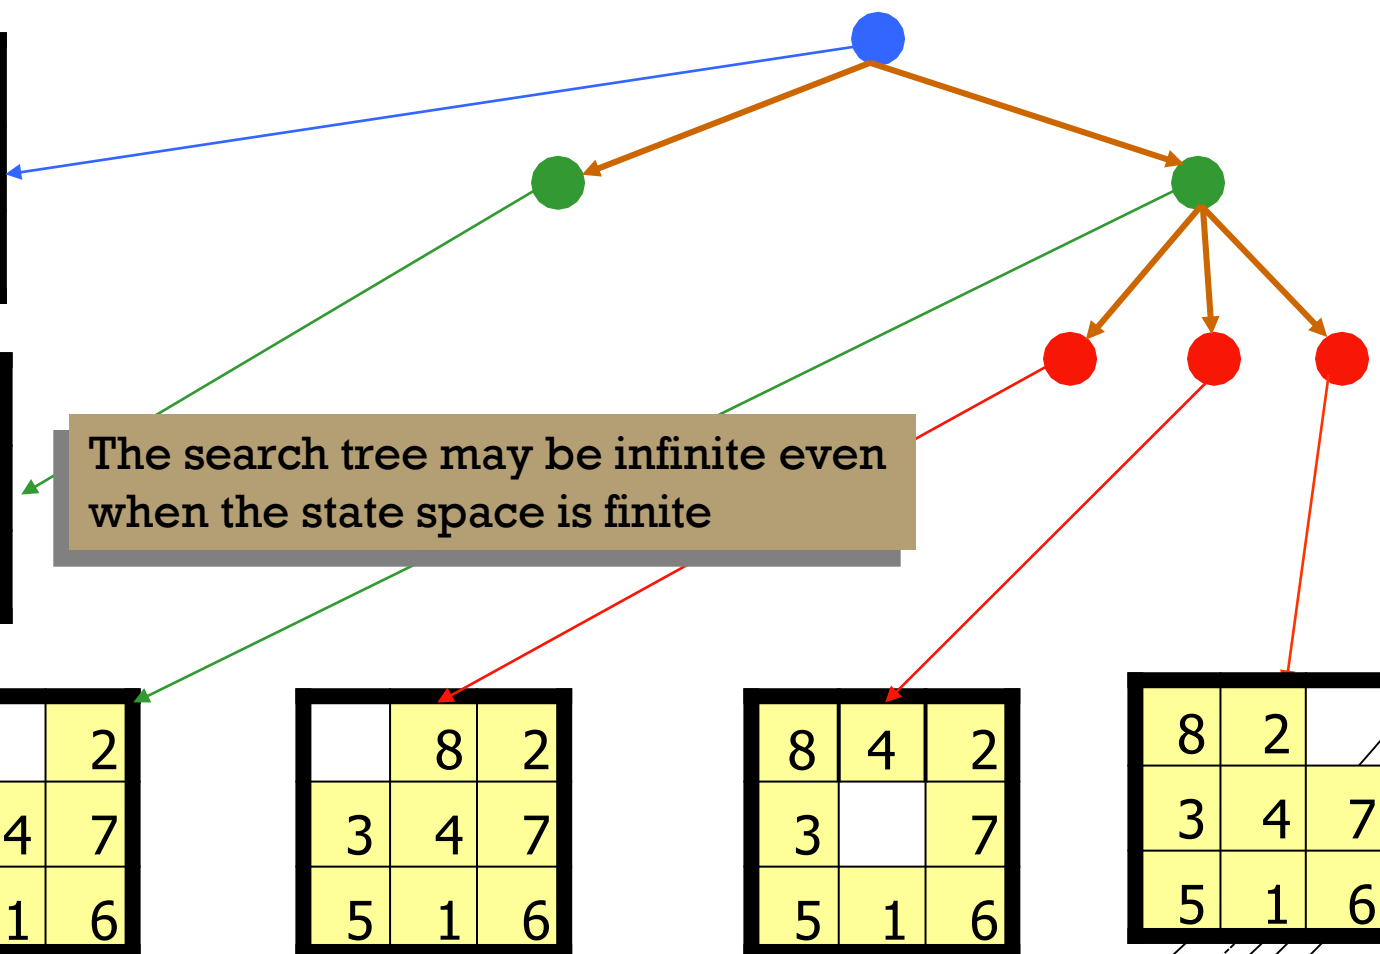
The search tree may be infinite even
when the state space is finite

8		2
3	4	7
5	1	6

	8	2
3	4	7
5	1	6

8	4	2
3		7
5	1	6

8	2	
3	4	7
5	1	6



Node Data Structure

- STATE
- PARENT
- ACTION
- COST

If a state is too large, it may be preferable to only represent the initial state and (re-)generate the other states when needed

Fringe

- Set of search nodes that have not been expanded yet
- Implemented as a **queue** FRINGE
 - INSERT(node,FRINGE)
 - REMOVE(FRINGE)
- The ordering of the nodes in FRINGE defines the search strategy

Search Algorithm

1. If GOAL?(initial-state) then return initial-state
2. INSERT(initial-node,FRINGE)
3. Repeat:
 - ◆ If FRINGE is empty then return failure
 - ◆ $n \leftarrow \text{REMOVE}(\text{FRINGE})$
 - ◆ $s \leftarrow \text{STATE}(n)$
 - ◆ For every state s' in SUCCESSORS(s)
 - Create a node n' as a successor of n
 - If GOAL?(s') then return path or goal state
 - INSERT(n' ,FRINGE)

Performance Measures

- **Completeness**
Is the algorithm guaranteed to find a solution when there is one?
- **Optimality**
Is this solution optimal?
- **Time complexity**
How long does it take?
- **Space complexity**
How much memory does it require?

Important Parameters

- Maximum number of successors of any state
→ branching factor **b** of the search tree
- Minimal length of a path in the state space between the initial and a goal state
→ depth **d** of the shallowest goal node in the search tree
- maximum length of a path in the state space
→ **m**

Important Remark

- Some problems formulated as search problems are NP-hard problems (e.g., (n^2-1) -puzzle).
- We cannot expect to solve such a problem in less than exponential time in the worst-case.
- But we can nevertheless strive to solve as many instances of the problem as possible.

Blind Strategies


- **Breadth-first**
 - Bidirectional
- **Depth-first**
 - Depth-limited
 - Iterative deepening
- **Uniform-Cost**

Step cost = 1



Two green arrows originate from the green box and point to the 'Bidirectional' sub-item under Breadth-first and the 'Depth-limited' sub-item under Depth-first.


Step cost = $c(\text{action})$
 $\geq \epsilon > 0$



A red arrow originates from the pink box and points to the 'Uniform-Cost' item in the list.

Breadth-First Strategy


New nodes are inserted **at the end** of
FRINGE



FRINGE = (1)

Breadth-First Strategy

New nodes are inserted **at the end** of
FRINGE




FRINGE = (2, 3)

Breadth-First Strategy

New nodes are inserted **at the end** of
FRINGE

FRINGE = (3, 4, 5)




Breadth-First Strategy

New nodes are inserted **at the end** of FRINGE


FRINGE = (4, 5, 6, 7)

[Node 7 (being the goal) could not actually be inserted into the fringe; shown here for clarity]



Bidirectional Strategy


2 fringe queues: FRINGE1 and FRINGE2



Time and space complexity = $O(b^{d/2}) \ll O(b^d)$

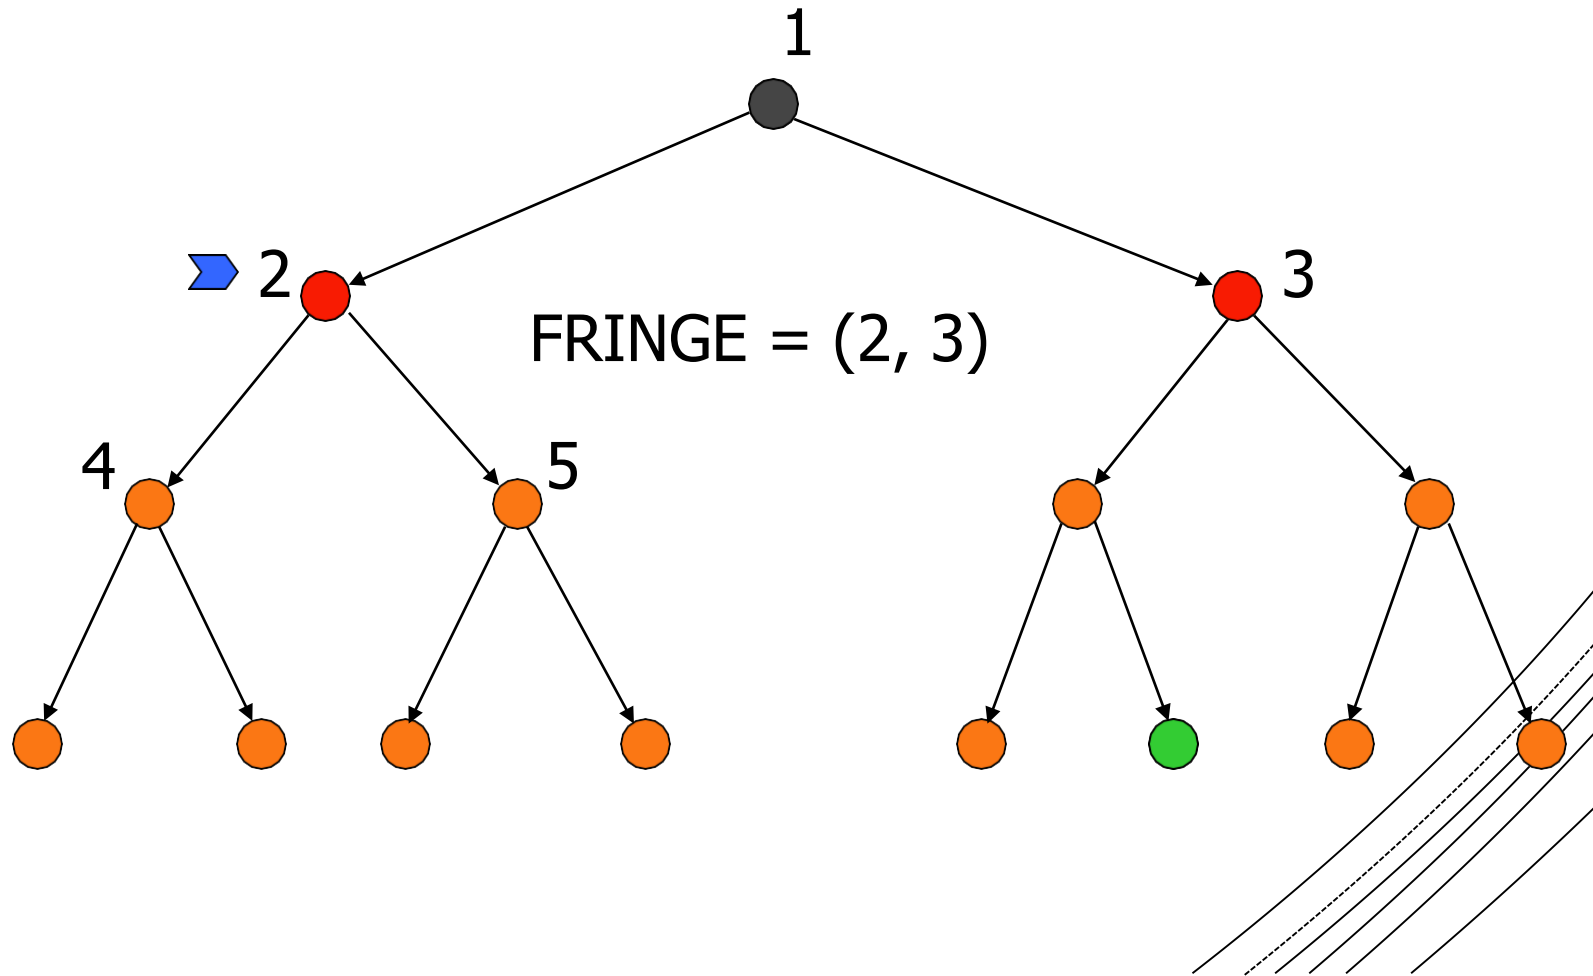
Depth-First Strategy

New nodes are inserted **at the front** of FRINGE




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front**
of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front** of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front** of FRINGE

Depth-First
Strategy




New nodes are inserted **at the front** of FRINGE

Depth-First
Strategy



New nodes are inserted **at the front** of FRINGE

Depth-First
Strategy





Depth-Limited Strategy

- Depth-first with **depth cutoff k** (maximal depth below which nodes are not expanded)
- Three possible outcomes:
 - Solution
 - Failure (no solution)
 - **Cutoff (no solution within cutoff)**

Uniform-Cost Strategy

- Each step has some cost $\geq \varepsilon > 0$.
- The cost of the path to each fringe node N is
$$g(N) = \sum \text{costs of all steps.}$$
- The goal is to generate a solution path of minimal cost.
- The queue FRINGE is sorted in increasing cost.



Modified Search Algorithm

1. INSERT(initial-node,FRINGE)
2. Repeat:
 - ◆ If FRINGE is empty then return **failure**
 - ◆ $n \leftarrow \text{REMOVE}(\text{FRINGE})$
 - ◆ $s \leftarrow \text{STATE}(n)$
 - ◆ If GOAL?(s) then return **path or goal state**
 - ◆ For every state s' in SUCCESSORS(s)
 - Create a node n' as a successor of n
 - INSERT(n' ,FRINGE)



So far ...

- Search tree \neq state space
- Search strategies: breadth-first, depth-first, and variants
- Evaluation of strategies: completeness, optimality, time and space complexity
- Avoiding repeated states
- Optimal search with variable step costs

Summary

- Search tree \neq state space
- Search strategies: breadth-first and depth-first
- Evaluation of strategies: completeness, optimality, time and space complexity

Acknowledgement

- Krishnaram Kenthapadi , CS121: Introduction to Artificial Intelligence, 2003.
- **AI Instructor's Resource Page,**
<http://aima.cs.berkeley.edu/instructors.html>